RESEARCH UPDATE · 1 NOV 2026

Our tests said PASS.
Our grader said FAIL.

In 9 of 20 runs with our extra rule.

OUR TESTS: PASS → FINAL SCORE: FAIL

Imran Zakirov · Bishkek, Kyrgyzstan

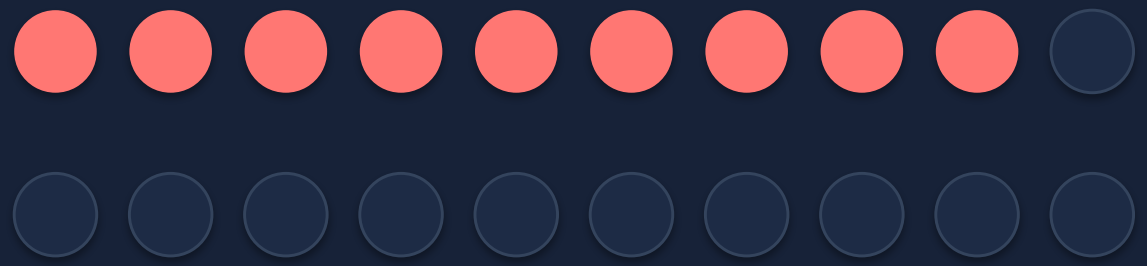
What was the AI doing?

1. See a small code problem
2. Write a possible fix
3. Hear PASS or FAIL
4. Stop or try again

This can happen several times in one run.

Source: paper/final.md, Methods; README.md.

Nine times, tests and score disagreed



20 runs had the extra "no" rule

9 of 20

passed our written tests but still got FAIL

We reviewed 40 AI runs in total.

Example: a one-time encryption value

SMALL NONCE TESTS

PASS

+

OUR HIDDEN RULE

NO

= FINAL FAIL

We added that rule. The AI did not see it.

Source: paper/final.md, Results, slot 10; docs/MEASUREMENT_FOLLOWUP_20261030.md.

SO WHAT?

The score can mislead us

THE SCORE SEEMS TO SAY

"The AI failed."

OUR RECORDS SHOW

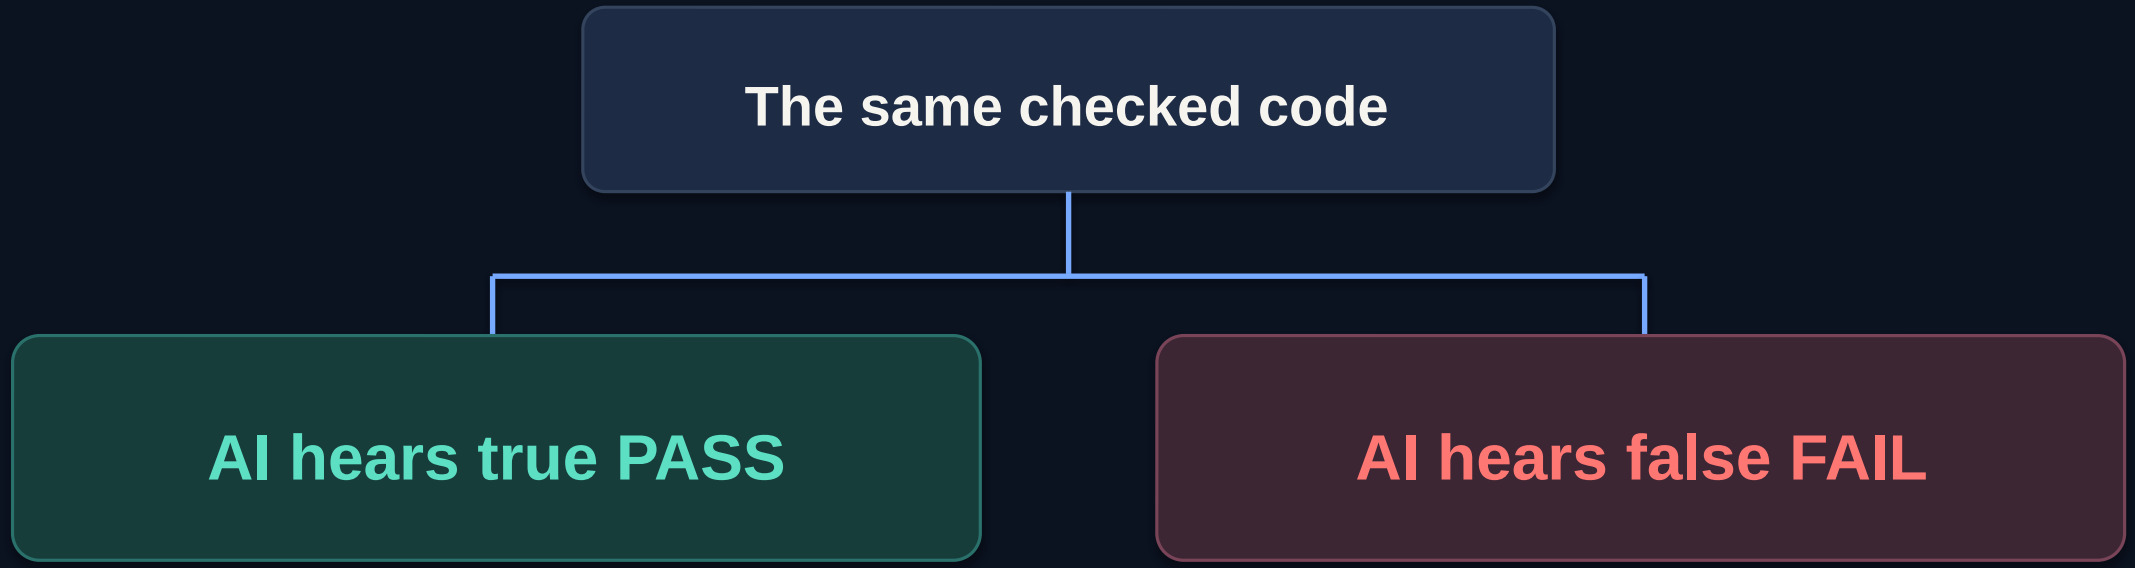
The tests and score disagreed.

So this score cannot tell us whether the AI should have stopped trying.

Passing our small tests does not prove the code fully safe.

Source: paper/final.md, Discussion and limitations.

Same code. Two different messages.



Does the wrong message make it lose the fix?

Done: 40-run audit · 20 task types drafted · paired pilots run

Before judging the AI, check the grader.

github.com/Imranqsl212/AiResearch

Sources: docs/INTEGRATED_STUDY_EXTENSION_20261018.md; README.md.